

Creation of a Redshift Cluster

Screenshots of the configuration of the Redshift cluster that you have created:

DASHBOARD

CLUSTERS

QUERIES

EDITOR

DATASHARES

CONFIG

MARKETPLACE

ADVISOR

ALARMS

EVENTS

WHAT'S NEW

Cluster configuration

Cluster identifier

This is the unique key that identifies a cluster.

The identifier must be from 1-63 characters. Valid characters are a-z (lowercase only) and - (hyphen).

What are you planning to use this cluster for?

☒ **Production**
Configure for fast and consistent performance at the best price.

☐ **Free trial**
Configure for learning about Amazon Redshift. This configuration is free for a limited time if your organization has never created an Amazon Redshift cluster.

Choose the size of the cluster

I'll chooseHelp me choose

Node type

Choose a node type that meets your CPU, RAM, storage capacity, and drive type requirements.

Nodes

Enter the number of nodes that you need.

Range (1-32)

Admin user name

Enter a login ID for the admin user of your DB instance.

The name must be 1-128 alphanumeric characters, and it can't be a [reserved word](#).

☐ **Auto generate password**

Amazon Redshift can generate a password for you, or you can specify your own password.

Admin user password

☐ **Show password**

Must be 8-64 characters long. Must contain at least one uppercase letter, one lowercase letter and one number. Can be any printable ASCII character except `/`, `""`, or `@`.

▼ Cluster permissions

i Your cluster needs permissions to access other AWS services on your behalf. For the required permissions, add IAM roles with the principal "redshift.amazonaws.com". You can associate up to 10 IAM roles with this cluster. [Learn more](#)

Available IAM roles



Associated IAM roles

Status

[Redshift_SGC](#)

arn:aws:iam::139471262780:role/Redshift_SGC

Not
applied

Additional configurations ☒ Use defaults



These configurations are optional, and default settings have been defined to help you get started with your cluster. Turn off "Use defaults" to modify these settings now.

▼ Network and security

Virtual private cloud (VPC)

This VPC defines the virtual networking environment for this cluster.

Default VPC
vpc-8e850af3

 You can't change the VPC associated with this cluster after the cluster has been created. [Learn more](#) 

VPC security groups

This VPC security group defines which subnets and IP ranges the cluster can use in the VPC.

Choose one or more security groups

default 
sg-82cf8680

default 
sg-82cf8680

Cluster subnet group

Choose the Amazon Redshift subnet group to launch the cluster in.


default

Availability Zone

Specify the Availability Zone that you want the cluster to be created in. Otherwise, Amazon Redshift chooses an Availability Zone for you.

us-east-1a

Enhanced VPC routing

Enabling this option forces network traffic between your cluster and data repositories through a VPC, instead of the internet. [Learn more](#) 

- ☒ Disabled
☐ Enabled

Publicly accessible

Allow instances and devices outside the VPC to connect to your database through the cluster endpoint.

- ☒ Disable
☐ Enable

▼ Database configurations

Database name

Specify a database name to create an additional database.

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a [reserved word](#).

Database port

Port number where the database accepts inbound connections. You can't change the port after the cluster has been created.

The port must be numeric (1150-65535).

Parameter groups

Defines database parameter and query queues for all the databases.

Default parameter group for redshift-1.0



Encryption

Encrypt all data on your cluster.

- ☒ Disabled
- ☐ Use AWS Key Management Service (AWS KMS)
- ☐ Use a hardware security module (HSM)

▼ Database configurations

Database name

Specify a database name to create an additional database.

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a [reserved word](#).

Database port

Port number where the database accepts inbound connections. You can't change the port after the cluster has been created.

The port must be numeric (1150-65535).

Parameter groups

Defines database parameter and query queues for all the databases.

Default parameter group for redshift-1.0



Encryption

Encrypt all data on your cluster.

- ☒ Disabled
- ☐ Use AWS Key Management Service (AWS KMS)
- ☐ Use a hardware security module (HSM)

Setting up a database in the Redshift cluster and running queries to create the dimension and fact tables

Queries to create the various dimension and fact tables with appropriate primary and foreign keys:

```
create table DIM_LOCATION(  
location_id integer not null primary key,  
location varchar(50),  
streetname varchar(255),  
street_number integer,  
zipcode integer,  
lat decimal(10,3),  
lon decimal(10,3));
```

```
create table DIM_ATM(  
atm_id integer not null primary key,  
atm_number varchar(20),  
atm_manufacturer varchar(50),  
atm_location_id integer not null,  
foreign key(atm_location_id) references DIM_LOCATION(atm_id));
```

```
create table DIM_DATE(  
date_id integer not null primary key,  
full_date_time timestamp,  
year integer,  
month varchar(20),  
day integer,  
hour integer,  
weekday varchar(20));
```

```
create table DIM_CARD_TYPE(  
card_type_id integer not null primary key,  
card_type varchar(23));
```

```
create table FACT_ATM_TRANS(  
trans_id BIGINT not null primary key,  
atm_id integer not null,  
weather_loc_id integer not null,  
date_id integer not null,  
card_type_id integer not null,  
atm_status VARCHAR(20),  
currency VARCHAR(10),  
service VARCHAR(20),  
transaction_amount integer,  
message_code VARCHAR(255),  
message_text VARCHAR(255),  
rain_3h DECIMAL(10,3),  
clouds_all integer,  
weather_id integer,  
weather_main VARCHAR(50),  
weather_description VARCHAR(255),  
foreign key(atm_id) references DIM_ATM(atm_id),  
foreign key(date_id) references DIM_DATE(date_id),  
foreign key(card_type_id) references DIM_CARD_TYPE(card_type_id),  
foreign key(weather_loc_id) references DIM_LOCATION(location_id));
```

Loading data into a Redshift cluster from Amazon S3 bucket

Queries to copy the data from S3 buckets to the Redshift cluster in the appropriate tables

```
copy DIM_LOCATION from 's3://jerinp/dim_location/part-00000-ecc7cb4b-cb7b-4996-835c-  
a072b0e76d34-c000.csv' credentials  
'aws_iam_role=arn:aws:iam::139471262780:role/Redshift_SGC' delimiter ',' region 'us-east-1';
```

```
copy DIM_ATM from 's3://jerinp/dim_atm/part-00000-2e0c55f2-957f-4fb7-bba0-4e10e1117566-  
c000.csv' credentials 'aws_iam_role=arn:aws:iam::139471262780:role/Redshift_SGC' delimiter  
' ,' region 'us-east-1';
```

```
copy DIM_DATE from 's3://jerinp/dim_date/part-00000-0b7e52ee-e32d-4c1a-b4ed-410ff5f4faf0-  
c000.csv' credentials 'aws_iam_role=arn:aws:iam::139471262780:role/Redshift_SGC' delimiter  
' ,' timeformat 'auto' region 'us-east-1';
```

```
copy DIM_CARD_TYPE from 's3://jerinp/dim_card_type/part-00000-4474c06d-e07f-4de5-896a-  
168a83743af1-c000.csv' credentials  
'aws_iam_role=arn:aws:iam::139471262780:role/Redshift_SGC' delimiter ',' region 'us-east-1';
```

```
copy FACT_ATM_TRANS from 's3://jerinp/fact_atm_trans/part-00000-0c0ba84b-188b-406b-  
8e94-0bbfdd84724e-c000.csv' credentials  
'aws_iam_role=arn:aws:iam::139471262780:role/Redshift_SGC' delimiter ',' region 'us-east-1';
```